

Modeling and Simulation in Python

Chapter 4

Copyright 2017 Allen Downey

License: Creative Commons Attribution 4.0 International

Lab Author @ Joseph Simone

```
# Configure Jupyter so figures appear in the notebook
%matplotlib inline

# Configure Jupyter to display the assigned value after an assignment
%config InteractiveShell.ast_node_interactivity='last_expr_or_assign'

# import functions from the modsim library
from modsim import *
```

Returning values

Here's a simple function that returns a value:

```
def add_five(x):
    return x + 5
```

And here's how we call it.

```
y = add_five(3)
```

8

If you run a function on the last line of a cell, Jupyter displays the result:

```
add_five(5)
```

10

But that can be a bad habit, because usually if you call a function and don't assign the result in a variable, the result gets discarded.

In the following example, Jupyter shows the second result, but the first result just disappears.

```
add_five(3)
add_five(5)
```

10

When you call a function that returns a variable, it is generally a good idea to assign the result to a variable.

```
y1 = add_five(3)
y2 = add_five(5)

print(y1, y2)
```

8 10

Exercise: Write a function called `make_state` that creates a `State` object with the state variables `olin=10` and `wellesley=2`, and then returns the new `State` object.

Write a line of code that calls `make_state` and assigns the result to a variable named `init`.

```
def make_state():
    state = State(olin=10, wellesley=2)
    return state
```

```
init = make_state()
```

values

olin

10

wellesley

2

Running simulations

Here's the code from the previous notebook.

```
def step(state, p1, p2):
    """Simulate one minute of time.

    state: bikeshare State object
    p1: probability of an Olin->Wellesley customer arrival
    p2: probability of a Wellesley->Olin customer arrival
    """
    if flip(p1):
        bike_to_wellesley(state)

    if flip(p2):
        bike_to_olin(state)

def bike_to_wellesley(state):
    """Move one bike from Olin to Wellesley.

    state: bikeshare State object
    """
    if state.olin == 0:
        state.olin_empty += 1
    return
```

```

state.olin -= 1
state.wellesley += 1

def bike_to_olin(state):
    """Move one bike from Wellesley to Olin.

    state: bikeshare State object
    """
    if state.wellesley == 0:
        state.wellesley_empty += 1
        return
    state.wellesley -= 1
    state.olin += 1

def decorate_bikeshare():
    """Add a title and label the axes."""
    decorate(title='Olin-Wellesley Bikeshare',
            xlabel='Time step (min)',
            ylabel='Number of bikes')

```

Here's a modified version of `run_simulation` that creates a `State` object, runs the simulation, and returns the `State` object.

```

def run_simulation(p1, p2, num_steps):
    """Simulate the given number of time steps.

    p1: probability of an Olin->Wellesley customer arrival
    p2: probability of a Wellesley->Olin customer arrival
    num_steps: number of time steps
    """
    state = State(olin=10, wellesley=2,
                  olin_empty=0, wellesley_empty=0)

    for i in range(num_steps):
        step(state, p1, p2)

    return state

```

Now `run_simulation` doesn't plot anything:

```
state = run_simulation(0.4, 0.2, 60)
```

values

olin

0

wellesley

12

olin_empty

5

```
wellesley_empty
```

```
0
```

But after the simulation, we can read the metrics from the `State` object.

```
state.olin_empty
```

```
5
```

Now we can run simulations with different values for the parameters. When `p1` is small, we probably don't run out of bikes at Olin.

```
state = run_simulation(0.2, 0.2, 60)
state.olin_empty
```

```
0
```

When `p1` is large, we probably do.

```
state = run_simulation(0.6, 0.2, 60)
state.olin_empty
```

```
13
```

More for loops

`linspace` creates a NumPy array of equally spaced numbers.

```
p1_array = linspace(0, 1, 5)
```

```
array([0. , 0.25, 0.5 , 0.75, 1.  ])
```

We can use an array in a `for` loop, like this:

```
for p1 in p1_array:
    print(p1)
```

```
0.0
0.25
0.5
0.75
1.0
```

This will come in handy in the next section.

`linspace` is defined in `modsim.py`. You can get the documentation using `help`.

```
help(linspace)
```

Help on function linspace in module modsim.modsim:

```
linspace(start, stop, num=50, **options)
    Returns an array of evenly-spaced values in the interval [start, stop].

    start: first value
    stop: last value
    num: number of values

    Also accepts the same keyword arguments as np.linspace. See
    https://docs.scipy.org/doc/numpy/reference/generated/numpy.linspace.html

    returns: array or Quantity
```

`linspace` is based on a NumPy function with the same name. Click [here](#) to read more about how to use it.

Exercise: Use `linspace` to make an array of 10 equally spaced numbers from 1 to 10 (including both).

```
zero_ten = linspace(0,10,10)
```

```
array([ 0.          ,  1.11111111,  2.22222222,  3.33333333,  4.44444444,
        5.55555556,  6.66666667,  7.77777778,  8.88888889, 10.          ])
```

```
for y in zero_ten:
    print(y)
```

```
0.0
1.1111111111111112
2.2222222222222223
3.3333333333333335
4.4444444444444445
5.5555555555555555
6.666666666666667
7.777777777777779
8.888888888888889
10.0
```

Exercise: The `modsim` library provides a related function called `linrange`. You can view the documentation by running the following cell:

```
help(linrange)
```

Help on function linrange in module modsim.modsim:

```
linrange(start=0, stop=None, step=1, endpoint=False, **options)
    Returns an array of evenly-spaced values in an interval.

    By default, the last value in the array is `stop-step`
    (at least approximately).
```

If you provide the keyword argument `endpoint=True`, the last value in the array is `stop`.

This function works best if the space between start and stop is divisible by step; otherwise the results might be surprising.

start: first value
stop: last value
step: space between values

returns: NumPy array

Use `linrange` to make an array of numbers from 1 to 11 with a step size of 2.

```
one_eleven = linrange(1,13,2)
```

```
array([ 1,  3,  5,  7,  9, 11], dtype=int32)
```

```
for z in one_eleven:  
    print(z)
```

```
1  
3  
5  
7  
9  
11
```

Sweeping parameters

`p1_array` contains a range of values for `p1`.

```
p2 = 0.2  
num_steps = 60  
p1_array = linspace(0, 1, 11)
```

```
array([0. , 0.1, 0.2, 0.3, 0.4, 0.5, 0.6, 0.7, 0.8, 0.9, 1. ])
```

The following loop runs a simulation for each value of `p1` in `p1_array`; after each simulation, it prints the number of unhappy customers at the Olin station:

```
for p1 in p1_array:  
    state = run_simulation(p1, p2, num_steps)  
    print(p1, state.olin_empty)
```

```
0.0 0  
0.1 0  
0.2 0  
0.30000000000000004 1  
0.4 10
```

```

0.5 15
0.6000000000000001 13
0.7000000000000001 28
0.8 28
0.9 36
1.0 39

```

Now we can do the same thing, but storing the results in a `SweepSeries` instead of printing them.

```

sweep = SweepSeries()

for p1 in p1_array:
    state = run_simulation(p1, p2, num_steps)
    sweep[p1] = state.olin_empty

```

And then we can plot the results.

```

plot(sweep, label='Olin')

decorate(title='Olin-Wellesley Bikeshare',
         xlabel='Arrival rate at Olin (p1 in customers/min)',
         ylabel='Number of unhappy customers')

```




Figure 1: png

Exercises

Exercise: Wrap this code in a function named `sweep_p1` that takes an array called `p1_array` as a parameter. It should create a new `SweepSeries`, run a simulation for each value of `p1` in `p1_array`, store the results in

the `SweepSeries`, and return the `SweepSeries`.

Use your function to plot the number of unhappy customers at Olin as a function of `p1`. Label the axes.

```
def sweep_p1():
    p2 = 0.2
    num_steps = 60
    p1_array = linspace(0, 1, 11)
    sweep = SweepSeries()

    for p1 in p1_array:
        state = run_simulation(p1, p2, num_steps)
        sweep[p1] = state.olin_empty
    return sweep
```

```
sweep_p1()
```

values

0.0

0

0.1

0

0.2

0

0.3

4

0.4

5

0.5

15

0.6

13

0.7

22

0.8

35

0.9

31

1.0

34


```

plot(sweep_p1(), label='Olin')

decorate(title='Olin-Wellesley Bikeshare',
         xlabel='Arrival rate at Olin (p1 in customers/min)',
         ylabel='Number of unhappy customers')

```




Figure 2: png

Exercise: Write a function called `sweep_p2` that runs simulations with `p1=0.5` and a range of values for `p2`. It should store the results in a `SweepSeries` and return the `SweepSeries`.

```

def sweep_p2():
    p1 = 0.5
    num_steps = 60
    p1_array = linspace(0, 1, 11)
    sweep = SweepSeries()

    for p2 in p1_array:
        state = run_simulation(p1, p2, num_steps)
        sweep[p2] = state.olin_empty
    return sweep

```

```
sweep_p2()
```

values

0.0

16

0.1
14
0.2
11
0.3
0
0.4
1
0.5
0
0.6
0
0.7
0
0.8
0
0.9
0
1.0
0

```
plot(sweep_p2(), label='Olin')

decorate(title='Olin-Wellesley Bikeshare',
          xlabel='Arrival rate at Olin (p1 in customers/min)',
          ylabel='Number of unhappy customers')
```

Optional Exercises

The following two exercises are a little more challenging. If you are comfortable with what you have learned so far, you should give them a try. If you feel like you have your hands full, you might want to skip them for now.

Exercise: Because our simulations are random, the results vary from one run to another, and the results of a parameter sweep tend to be noisy. We can get a clearer picture of the relationship between a parameter and a metric by running multiple simulations with the same parameter and taking the average of the results.

Write a function called `run_multiple_simulations` that takes as parameters `p1`, `p2`, `num_steps`, and `num_runs`.

`num_runs` specifies how many times it should call `run_simulation`.

After each run, it should store the total number of unhappy customers (at Olin or Wellesley) in a `TimeSeries`. At the end, it should return the `TimeSeries`.

Test your function with parameters




Figure 3: png

```
p1 = 0.3
p2 = 0.3
num_steps = 60
num_runs = 10
```

Display the resulting `TimeSeries` and use the `mean` function provided by the `TimeSeries` object to compute the average number of unhappy customers (see Section 2.7).

```
def run_multiple_simulations(p1,p2,num_steps,num_runs):
    result = TimeSeries()
    for i in range(num_runs):
        state = run_simulation(p1, p2, num_steps)
        result[i] = state.olin
    return result
```

```
run_time = run_multiple_simulations(0.3,0.3,60,10)
```

values

```
0
10
1
12
2
8
```

3
12
4
5
5
1
6
8
7
6
8
7
9
7

```
plot(run_time, label='Olin')  
  
decorate(title='Olin-Wellesley Bikeshare',  
          xlabel='Arrival rate at Olin (p1 in customers/min)',  
          ylabel='Number of unhappy customers')
```




Figure 4: png

```
run_time.mean()
```

7.6

Exercise: Continuing the previous exercise, use `run_multiple_simulations` to run simulations with a range of values for `p1` and

```
p2 = 0.3
num_steps = 60
num_runs = 20
```

Store the results in a `SweepSeries`, then plot the average number of unhappy customers as a function of `p1`. Label the axes.

What value of `p1` minimizes the average number of unhappy customers?

```
def run_multiple_simulations2(p2,num_steps,num_runs):
    result = SweepSeries()
    p1_array = linspace(0, 1, 11)
    for i in range(num_runs):
        for p1 in p1_array:
            state = run_simulation(p1, p2, num_steps)
            result[p1] = state.olin
    return result
```

```
run_time_2 = run_multiple_simulations2(0.3,60,20)
```

values

0.0

12

0.1

12

0.2

11

0.3

2

0.4

0

0.5

1

0.6

0

0.7

1

0.8
1
0.9
0
1.0
1

```
plot(run_time_2, label='Olin')  
  
decorate(title='Olin-Wellesley Bikeshare',  
          xlabel='Arrival rate at Olin (p1 in customers/min)',  
          ylabel='Number of unhappy customers')
```




Figure 5: png

```
run_time_2.mean()
```

3.727272727272727